



Choosing the Smallest LLM That Won't Completely Fail You

ContainerDays London '26

\$> *what are we going to see today?*

01. SLMs
02. Run local models (with Open Source!)
03. Benchmark models (with Open Source!)



\$> tell me everything you know about me



Manuel de la Peña

Staff Software Engineer

Docker

- 2022 → AtomicJar, acquired by Docker
- 2019 → Elastic: Eng. Productivity (Observability)
- 2011 → Liferay: Core Engineer > QA Tech Lead (Cloud)
- 2008 → Consultant
- 2005 → Public Admin in Spain
- Hitting keyboards since 1994

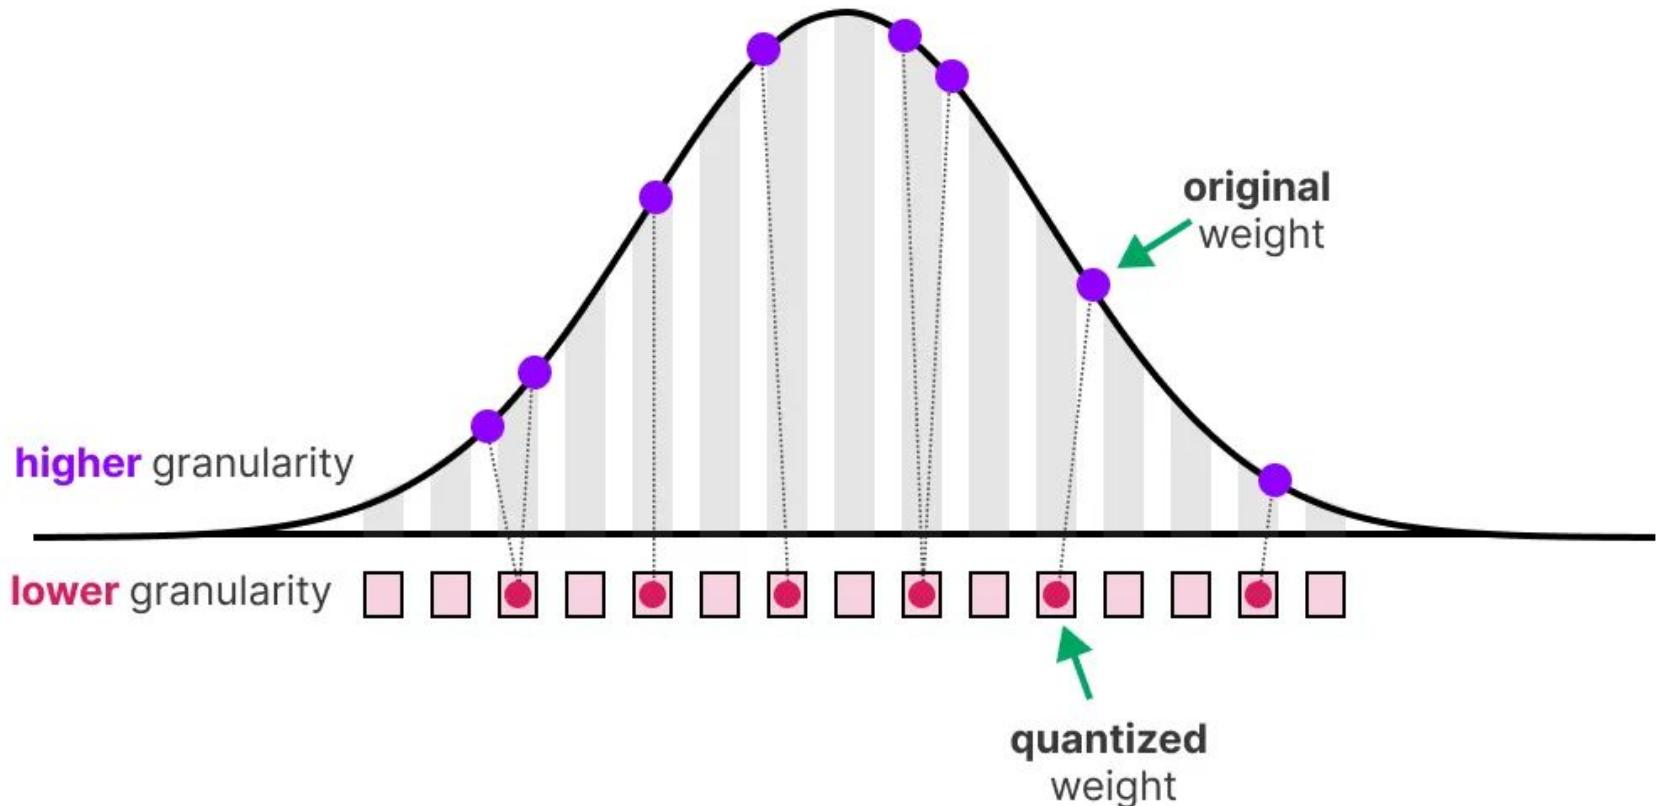


\$> what are Small Language Models (SLMs)?

- Mostly Open Source models
- They can fit in your host's GPUs
 - ◆ 2-8 Gb of disk (and GPUs)
- Not producing as rich responses as the Big Players' models
- Ideal for experimentation, or for specialised tasks
- Multiple flavours (by **quantization**)
 - ◆ Reducing how the weights are stored in floating point numbers
 - FP32: 100% of the size
 - FP16: 50% of the size



\$> draw a graph representing quantization



\$> \textbf{\textit{what are the trade-offs for using quantization?}}

Quantization comes with trade-offs:

- ◆ Slightly worse reasoning
- ◆ More hallucinations at higher temperatures
- ◆ Sensitive tasks (math, code) degrade sooner



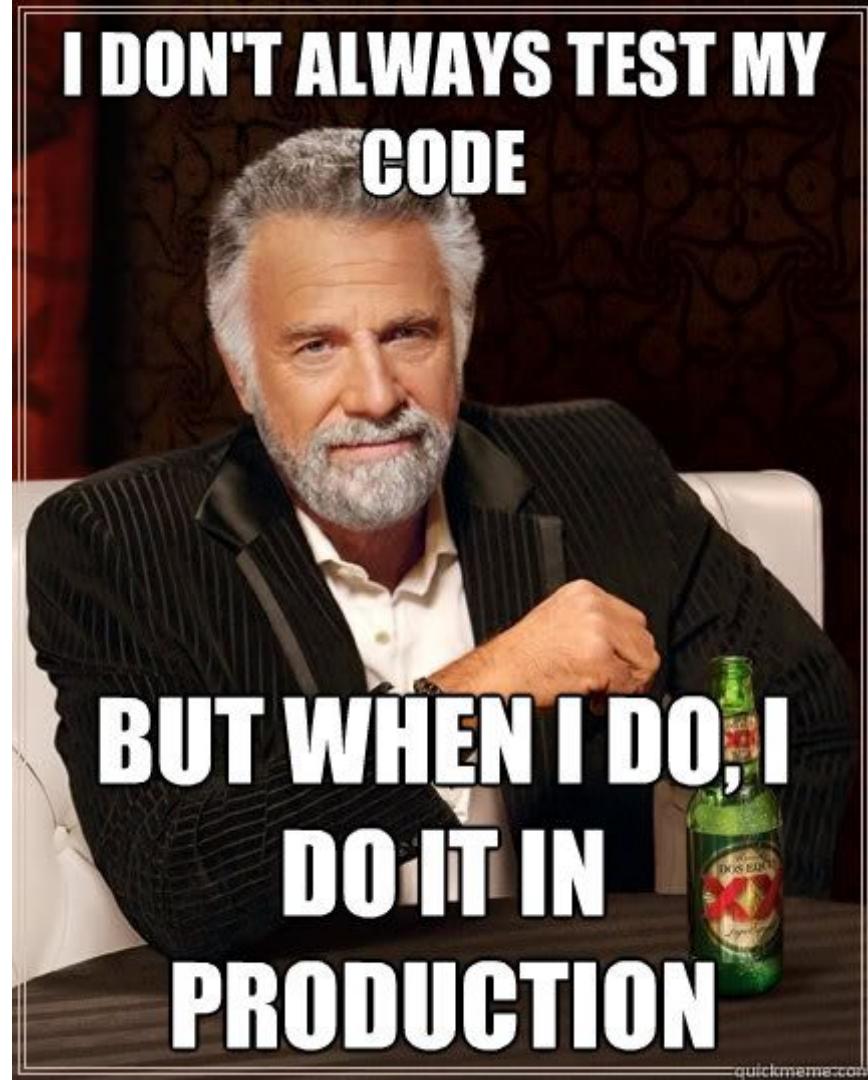
\$> Reality is ...

You may still use the LARGE models

- ◆ Way RICHER responses and outputs

Would you use the production models to test every change while building the app?

- Local development
- CI



\$> /stats

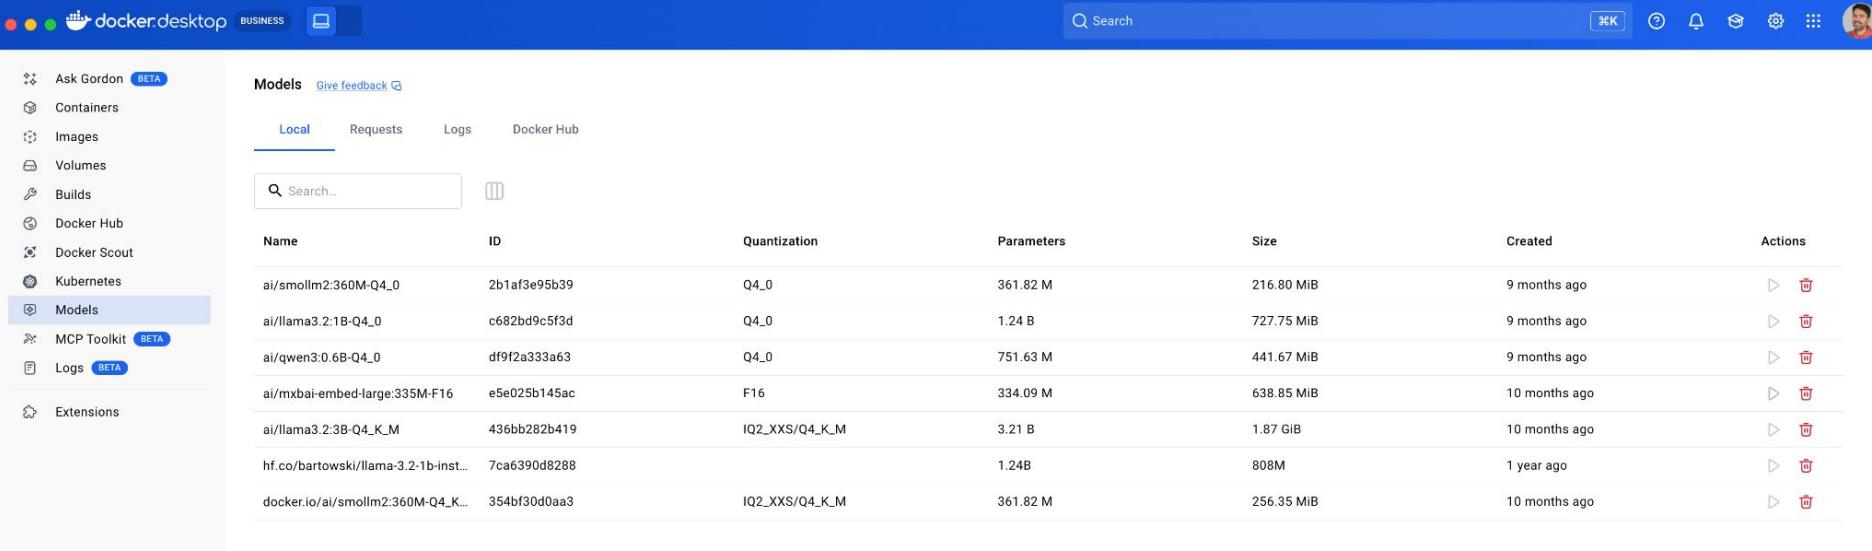




THISISSELLZ

\$> how can I run models locally with Open Source?

- Llama.cpp, [ardanlabs/kronk](#) (Go library), [Ollama...](#)
- [Docker Model Runner](#) (Linux, MacOS and Windows)



The screenshot shows the Docker Desktop interface with the 'Models' section selected. The sidebar on the left has 'Models' highlighted with a blue background. The main area displays a table of local models with columns for Name, ID, Quantization, Parameters, Size, Created, and Actions.

Name	ID	Quantization	Parameters	Size	Created	Actions
ai/smollm2:360M-Q4_0	2b1af3e95b39	Q4_0	361.82 M	216.80 MiB	9 months ago	View Edit
ai/llama3.2:1B-Q4_0	c682bd9c5f3d	Q4_0	1.24 B	727.75 MiB	9 months ago	View Edit
ai/qwen3.0:6B-Q4_0	df9f2a333a63	Q4_0	751.63 M	441.67 MiB	9 months ago	View Edit
ai/mxbai-embed-large:335M-F16	e5e025b145ac	F16	334.09 M	638.85 MiB	10 months ago	View Edit
ai/llama3.2:3B-Q4_K_M	436bb282b419	IQ2_XXS/Q4_K_M	3.21 B	1.87 GiB	10 months ago	View Edit
hf.co/bartowski/llama-3.2-1b-inst...	7ca6390d8288		1.24B	808M	1 year ago	View Edit
docker.io/ai/smollm2:360M-Q4_K_M	354bf30d0aa3	IQ2_XXS/Q4_K_M	361.82 M	256.35 MiB	10 months ago	View Edit



\$> how and where can I get the models?

→ [Docker Hub: AI Models](#)

- ◆ Models as OCI artifacts → **docker model pull ai/llama3.2:3B-Q4_0**
- ◆ TIP: to download the recommended model → **docker model pull gpt-oss**

→ [Hugging Face](#)

```
● ● ●
docker model pull ai/llama3.2:3B-Q4_0
1e817ea420d7: Pull complete [=====>] 1.917GB/1.917GB
Model pulled successfully
```



\$> can I run models with Docker Model Runner?



```
docker model run ai/llama3.2:3B-Q4_0
> do you like testcontainers-go?
Testcontainers-go is a fantastic tool for testing Kubernetes applications in Go. It simplifies the process of setting up ephemeral containers for testing, allowing you to write cleaner and more efficient tests.
```

Here are some reasons why I think Testcontainers-go is great:

1. **Easy setup**: Testcontainers-go makes it easy to set up a container for testing, reducing the amount of boilerplate code you need to write.
2. **Ephemeral containers**: By default, Testcontainers-go creates ephemeral containers, which means they are automatically deleted after the test is finished. This reduces the overhead of running long-running containers.
3. **Customizable**: You can customize the container image, command, and environment variables to suit your testing needs.
4. **Multi-environment support**: Testcontainers-go supports multiple environments, including local, Docker Desktop, and Kubernetes clusters.
5. **Integration with testing frameworks**: Testcontainers-go integrates well with popular testing frameworks like GoTest and Testify.

However, as with any library, it's essential to consider the following:

1. **Performance**: Running containers for testing can impact performance, especially if you're running multiple tests concurrently.
2. **Resource usage**: Testcontainers-go uses system resources, so be mindful of the impact on your test environment.
3. **Cleanup**: Make sure to properly clean up after each test to avoid leaving resources idle.

Overall, Testcontainers-go is a valuable tool for writing efficient, reliable, and scalable tests for Kubernetes applications in Go.
> Send a message (/? for help)



\$> explain Docker Hub models names

Models in Docker Hub use the recommended quantization to not lose accuracy and reduce size.

I.e. llama3.2:3B-Q4_K_M

- Number of **parameters** (3B/3b for 3 billions)
- **Quantization** (Q4/q4 for 4-bit quantization → using 4 bits instead of 32)
- **K-quantization** (from llama.cpp) uses groups quantization for better accuracy
- **Medium variant of K**, balanced choice between speed, memory and accuracy

- So “**q4_k_m**” is usually a sweet spot for local and edge use.





Benchmark models with Open Source

\$> what pieces do you find in a regular AI software?

When building AI software, we usually need to define the following:

- System prompts for the models to define initial behaviors
- Model parameters (Temperature, Top_P, Top_K...)
- The model itself (Reasoning capabilities, Using Tools...)
- Response formats (Natural language, JSON, ...)

How do you know that your choices are the best for your use case?

- ◆ MMLU benchmark (Massive Multitask Language Understanding)
- ◆ Benchmarks by third-parties

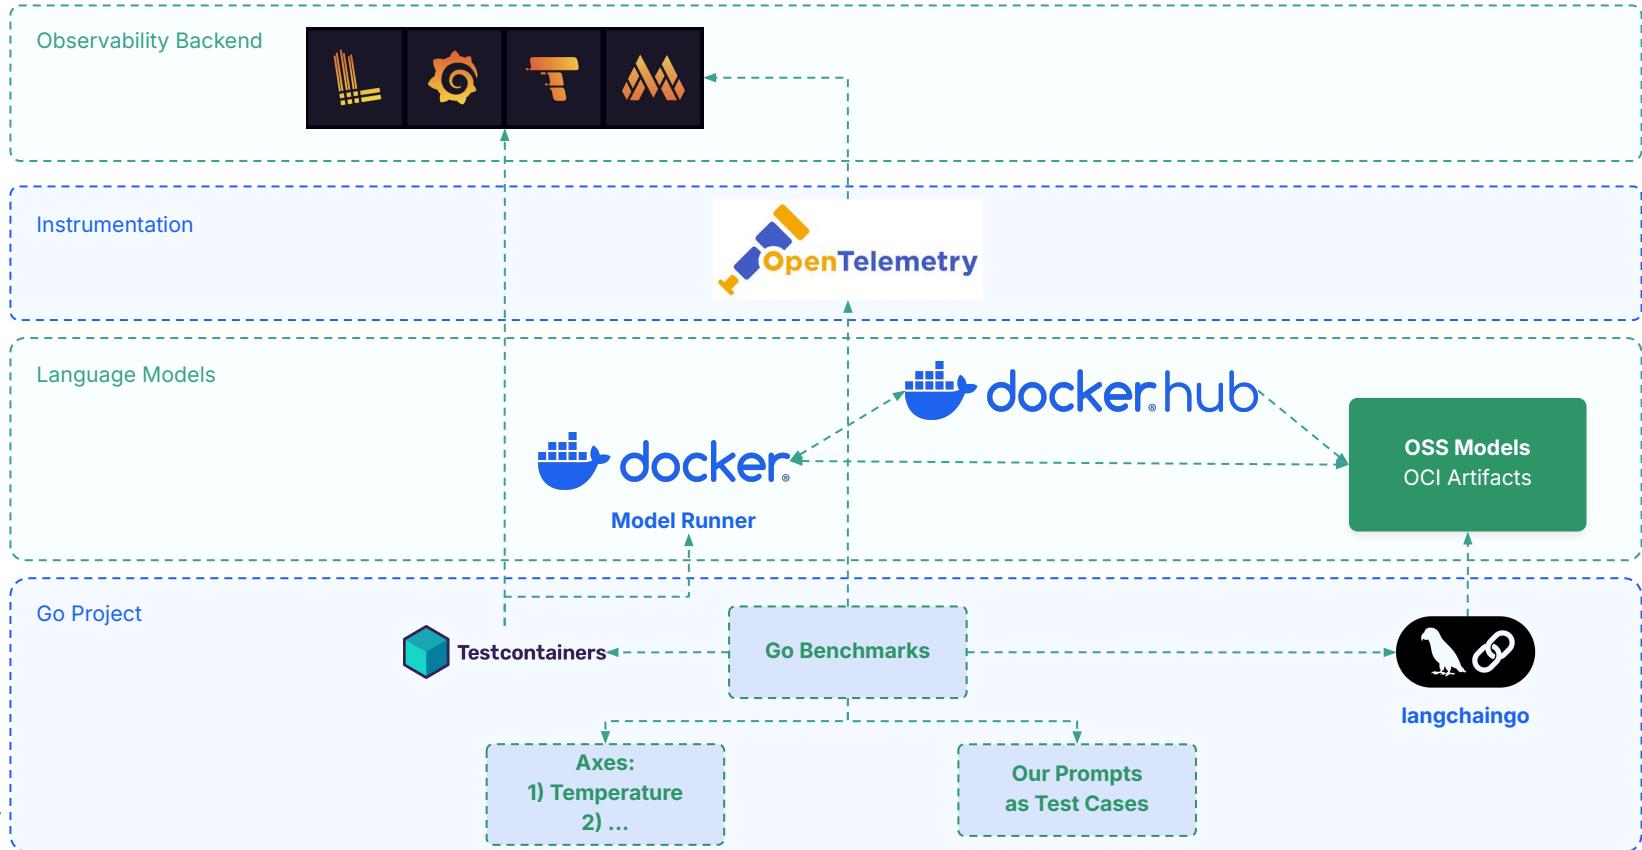


\$> list OSS-only tools I use to benchmark models

- Golang Benchmarks from the testing package
- [Docker Model Runner](#) (DMR) to load the OSS models
 - ◆ gpt-5.1
 - ◆ ai/llama3.2:3B-Q4_K_M
 - ◆ ai/llama3.2:1B-Q4_0
 - ◆ ai/qwen3:0.6B-Q4_0
 - ◆ hf.co/bartowski/Llama-3.2-1B-Instruct-GGUF
- OpenTelemetry (OTEL) for contributing metrics, traces and logs
- Grafana LGTM as observability backend
- [Langchain Go](#) for talking to the models in Go
- Testcontainers Go for starting containers (using predefined modules!)
 - ◆ Grafana, Python (for code execution)
 - ◆ Socat to talk to DMR through [model-runner.docker.internal:80](#)



\$> draw an architectural diagram



\$> design the benchmarks ($8 \times 5 \times 5 = 200$ scenarios)

→ Axe-0: 8 task types

- ◆ **Code Explanation** of a function implementing Fibonacci.
- ◆ **Mathematical operations** (what is the sum of all numbers between 1 and 100, both inclusive?=)
- ◆ **Creative Writing**: Write a hilarious joke about the Fibonacci sequence.
- ◆ **Factual Question**: What was the significance of Toledo, Spain during the medieval period, particularly regarding the translation movement?
- ◆ **Code Generation**: Write a Go function that calculates the Fibonacci sequence using recursion.
- ◆ **Calculator Reasoning**: Break down complex arithmetic into multiple calculator calls (
- ◆ **Code Validation**: Generate Python code for Fibonacci, and execute it via a code executor tool
- ◆ **Api Data Retrieval**: Use Fetch tool to get GitHub repository data and summarizes the key details

→ Axe-1: 5 Temperatures → 0.1, 0.3, 0.5, 0.7, 0.9

→ Axe-2: 5 OSS Models



\$> what metrics could we collect from a model?

→ Metrics:

- ◆ Latency Percentiles: overall response time
- ◆ Latency Histogram: response time distribution
- ◆ Prompt Evaluation Time: Time to first token (TTFT)
- ◆ Tokens per Operation: Token usage and verbosity
- ◆ Tokens per Second: Generation throughput
- ◆ Success Rate: Model reliability
- ◆ GPU Utilization: Hardware efficiency
- ◆ GPU Memory: Consumption
- ◆ Score per Operation: Quality metrics provided by an Evaluator agent





Code Demo



GitHub repository

\$> summarise what we have learn today

01. OSS is cool
02. SLM models and quantization
03. Using OSS tools to benchmark models
04. Plot the benchmarks into an OTEL backend





\$> *You are absolutely cool!*
Thanks

ContainerDays London '26



Annexes



Massively Large models Vs Open Source models

(Massively) Large Language Models

Big players have huge models with a humongous number of parameters.

Real numbers are not disclosed

→ OpenAI

- ◆ GPT-5.2 → trillions?
- ◆ GPT 4o → 100 billions?

→ Anthropic:

- ◆ Claude 4 Sonnet → 300-500 billions?
- ◆ Opus → 1-2 trillions?

→ Google:

- ◆ Gemini 3 → 1-3 trillions?



Open Source Language Models

Ideal to run locally or in your own infrastructure, ideally on GPUs

Model variant	Parameters	Quantization	Context window	VRAM ¹	Size
ai/llama3.2:latest	3B	IQ2_XXS/Q4_K_M	131K tokens	2.77 GiB	1.87 GB
ai/llama3.2:3B-Q4_K_M					
ai/llama3.2:1B-Q4_0	1B	Q4_0	131K tokens	1.35 GiB	727.75 MB
ai/llama3.2:1B-Q8_0	1B	Q8_0	131K tokens	1.87 GiB	1.22 GB
ai/llama3.2:1B-F16	1B	F16	131K tokens	2.95 GiB	2.30 GB
ai/llama3.2:3B-Q4_0	3B	Q4_0	131K tokens	2.68 GiB	1.78 GB
ai/llama3.2:3B-Q4_K_M	3B	IQ2_XXS/Q4_K_M	131K tokens	2.77 GiB	1.87 GB
ai/llama3.2:3B-F16	3B	F16	131K tokens	6.89 GiB	5.98 GB



... AND size matters!

The bigger the model, the more expensive using it

GPT-5.2

The best model for coding and
agentic tasks across industries

Price

Input:
\$1.750 / 1M tokens

Cached input:
\$0.175 / 1M tokens

Output:
\$14.000 / 1M tokens

GPT-5.2 pro

The smartest and most
precise model

Price

Input:
\$21.00 / 1M tokens

Cached input:
-

Output:
\$168.00 / 1M tokens

GPT-5 mini

A faster, cheaper version of GPT-5
for well-defined tasks

Price

Input:
\$0.250 / 1M tokens

Cached input:
\$0.025 / 1M tokens

Output:
\$2.000 / 1M tokens



Source: <https://openai.com/api/pricing/>

... AND size matters!

The bigger the model, the more expensive using it

Model	Base Input Tokens	5m Cache Writes	1h Cache Writes	Cache Hits & Refreshes	Output Tokens
Claude Opus 4.5	\$5 / MTok	\$6.25 / MTok	\$10 / MTok	\$0.50 / MTok	\$25 / MTok
Claude Opus 4.1	\$15 / MTok	\$18.75 / MTok	\$30 / MTok	\$1.50 / MTok	\$75 / MTok
Claude Opus 4	\$15 / MTok	\$18.75 / MTok	\$30 / MTok	\$1.50 / MTok	\$75 / MTok
Claude Sonnet 4.5	\$3 / MTok	\$3.75 / MTok	\$6 / MTok	\$0.30 / MTok	\$15 / MTok
Claude Sonnet 4	\$3 / MTok	\$3.75 / MTok	\$6 / MTok	\$0.30 / MTok	\$15 / MTok
Claude Sonnet 3.7 <u>(deprecated)</u>	\$3 / MTok	\$3.75 / MTok	\$6 / MTok	\$0.30 / MTok	\$15 / MTok
Claude Haiku 4.5	\$1 / MTok	\$1.25 / MTok	\$2 / MTok	\$0.10 / MTok	\$5 / MTok
Claude Haiku 3.5	\$0.80 / MTok	\$1 / MTok	\$1.6 / MTok	\$0.08 / MTok	\$4 / MTok
Claude Opus 3 <u>(deprecated)</u>	\$15 / MTok	\$18.75 / MTok	\$30 / MTok	\$1.50 / MTok	\$75 / MTok
Claude Haiku 3	\$0.25 / MTok	\$0.30 / MTok	\$0.50 / MTok	\$0.03 / MTok	\$1.25 / MTok



Source: <https://platform.claude.com/docs/en/about-claude/pricing>

... AND size matters!

The bigger the model, the more expensive using it

Standard	Batch	
	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Not available	\$2.00, prompts <= 200k tokens \$4.00, prompts > 200k
Output price (including thinking tokens)	Not available	\$12.00, prompts <= 200k tokens \$18.00, prompts > 200k
Context caching price	Not available	\$0.20, prompts <= 200k tokens \$0.40, prompts > 200k \$4.50 / 1,000,000 tokens per hour (storage price)
Grounding with Google Search*	Not available	5,000 prompts per month (free), then (Coming soon**) \$14 / 1,000 search queries
Grounding with Google Maps	Not available	Not available
Used to improve our products	Yes	No



Source: <https://ai.google.dev/gemini-api/docs/pricing>